

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

word2vec

node2vec

Summary

# Embeddings

Stephen Scott

(Adapted from Haluk Dogan)

[sscott2@unl.edu](mailto:sscott2@unl.edu)

# Introduction

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectors

word2vec

node2vec

Summary

- To apply recurrent and transformer architectures to text (e.g., NLM), need numeric representation of words
  - The “Embedding lookup” block
  - Want words with similar meanings to cluster together
- Where does the embedding come from?
  - Could train it along with the rest of the network
  - Or, could use “off-the-shelf” embedding
    - E.g., word2vec or GloVe
- Embeddings not limited to words: E.g., biological sequences, graphs, ...
  - Graphs: node2vec
- The xxxx2vec approach focuses on training embeddings based on **context** of what words often appear near it
  - Different “context” than, e.g., BERT, which has separate embedded vector for each usage

# Outline

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

word2vec

node2vec

Summary

- Vector semantics
- Representing words as vectors
- word2vec
  - Architectures
  - Training
  - Semantics of embedding
- node2vec
- Summary

# Vector Semantics

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

word2vec

node2vec



Summary

- Want to numerically represent words such that words sharing lexical semantics tend to cluster together
  - A trained model that predicts positive when the word “vanish” appears ought to respond similarly to the word “disappear”
- Will define representation that groups words together when they appear in similar **contexts**

# Vector Semantics

[Embeddings](#)[Stephen Scott](#)[Introduction](#)[Vector  
semantics](#)[Words and  
Vectors](#)[word2vec](#)[node2vec](#)[Summary](#)

## Two-dimensional (t-SNE) projection of 60-dimensional embedding



# Words and Vectors

## Term-Document Matrix

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
VectorsTerm-Document  
Matrix

Words as Vectors

Cosine Similarity

TF-IDF

word2vec

node2vec

Summary

- Let  $D$  be set of documents based on vocabulary  $V$
- A **term-document matrix** is a  $|V| \times |D|$  matrix where column  $d$  is document  $d$ 's bag-of-words representation (word counts)

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	0	7	13
good	114	80	62	89
fool	36	58	1	4
wit	20	15	2	3

# Words and Vectors

## Term-Document Matrix

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

Term-Document  
Matrix

Words as Vectors

Cosine Similarity


TF-IDF

word2vec

node2vec

Summary

### Two-dimensional projection



Comedies cluster, as do non-comedies

# Words and Vectors

## Representing Words as Vectors

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

Term-Document  
Matrix

Words as Vectors

Cosine Similarity  
TF-IDF

word2vec

node2vec

Summary

- Can use the rows of the T-D matrix to represent words
- Context for a word is the set of documents that a word appears in, and how often
- Expect similar words to have similar representations
- Can also use a  $|V| \times |V|$  **word-word** (or **term-context**) matrix
  - Cell  $(i,j)$  holds the number of times words  $i$  and  $j$  appear in the same **context** (same document, same window, etc.)

# Words and Vectors

## Representing Words as Vectors

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
VectorsTerm-Document  
Matrix

Words as Vectors

Cosine Similarity


TF-IDF

word2vec

node2vec

Summary

	aardvark	...	computer	data	result	pie	sugar	...
cherry	0	...	2	8	9	442	25	
strawberry	0	...	0	0	1	60	19	
digital	0	...	1670	1683	85	5	4	
information	0	...	3325	3982	378	5	13	



# Cosine Similarity

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
VectorsTerm-Document  
Matrix

Words as Vectors

Cosine Similarity

TF-IDF

word2vec


node2vec

Summary

- Can compute similarity of words (or documents)
- Common approach: The **cosine** of the angle between vectors  $v$  and  $w$ :

$$\frac{v \cdot w}{\|v\| \|w\|} = \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$

- I.e., **normalize** each vector by dividing by its length, then take the dot product of unit-length vectors
- Similarity = 1 if vectors identical, 0 if orthogonal (never  $< 0$  since all vectors nonnegative)



# Term Frequency-Inverse Document Frequency

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

Term-Document  
Matrix

Words as Vectors

Cosine Similarity

TF-IDF

word2vec

node2vec

Summary

- Want to be careful about how we define context
- Common words such as “the” appear frequently in any corpus
- Can artificially inflate similarity measures
- TF-IDF** weights terms not only by frequency, but also **reduces** a weight if a term appears in many documents

# Term Frequency-Inverse Document Frequency

## Definition

[Embeddings](#)[Stephen Scott](#)[Introduction](#)[Vector  
semantics](#)[Words and  
Vectors](#)[Term-Document  
Matrix](#)[Words as Vectors](#)[Cosine Similarity](#)[TF-IDF](#)[word2vec](#)[node2vec](#)[Summary](#)

- **Term frequency** counts frequency of term  $t$  in document  $d$ , with a 1 pseudocount to avoid log of zero:

$$\text{tf}_{t,d} = \log_{10}(\text{count}(t, d) + 1)$$

- Log is used since, e.g., 100 occurrences is not much different than 200
- **Inverse document frequency** decreases as the number of documents containing  $t$  increases:

$$\text{idf}_t = \log_{10} \left( \frac{N}{\text{df}_t} \right)$$

$N$  = total number of documents,  $\text{df}_t$  = number of documents containing  $t$

- Weight of  $t, d$  is then  $w_{t,d} = (\text{tf}_{t,d})(\text{idf}_t)$

# Term Frequency-Inverse Document Frequency

## Use

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

Term-Document  
Matrix

Words as Vectors

Cosine Similarity

TF-IDF

word2vec

node2vec

Summary

- Now, when we compute word-word matrix, weight counts by tf-idf
- Then extract vector for term  $t$  as usual
- Can then normalize and compute cosine similarity

# Word2vec (Mikolov et al.)

## Dense Vector Representations

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

word2vec

Intuition

Training

Visualization

Semantics

Miscellany

node2vec

Summary

- Vectors based on tf-idf are typically long ( $|V|$  order of  $10^4$ – $10^5$ ) and sparse (mostly 0 entries)
- Alternative: Vectors that are short (say, 100 dimensions) and dense
  - Fewer model parameters to learn
  - Capture synonymy: In a sparse representation, “car” and “automobile” are distinct dimensions, so vectors might be distant
  - Work better in practice
- Examples
  - Word2vec (<https://code.google.com/archive/p/word2vec/>)
  - Fasttext (<http://www.fasttext.cc/>)
  - GloVe (<http://nlp.stanford.edu/projects/glove/>)

# Word2vec (Mikolov et al.)

## Intuition

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectors

word2vec

Intuition

Training

Visualization

Semantics

Miscellany

node2vec

Summary

- Given a word  $w$ 
  - Rather than **count** how often  $w$  appears near word “apricot”
  - **Predict** how likely  $w$  is to appear near “apricot”
- **Equivalent view:** Train as a variation of autoencoding
  - But rather than mapping a word to itself, learn to map between a word and its **context**
  - Context-to-word: **Continuous bag-of-words** (CBOW)
  - Word-to-context: **Skip-gram**
- Can train this with a neural network
  - Can use supervised training, since any word appearing near “apricot” in corpus is a positive example
  - Will also use **negative sampling** in training
- Model’s weights can be used as word embeddings

# Word2vec (Mikolov et al.)

## Architectures

Embeddings


Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectorsword2vec  
Intuition  
TrainingVisualization  
Semantics  
Miscellany

node2vec

Summary



- CBOW: Predict current word  $w(t)$  based on context
- Skip-gram: Predict context based on  $w(t)$
- One-hot input, hidden linear activation, softmax output

# Word2vec (Mikolov et al.)

## CBOW

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectorsword2vec  
Intuition

Training


Visualization

Semantics

Miscellany

node2vec

Summary



- Use  $i$ th column of  $W'$  as embedding

- $N = \text{vocabulary size}, d = \text{embedding dimension}$
- $N \times d$  matrix  $W$  is shared weights from input to hidden
- $d \times N$  matrix  $W'$  is weights from hidden to output
- When one-hot context vectors  $x_{t-2}, x_{t-1}, \dots, x_{t+2}$  input, corresponding rows from  $W$  are summed to  $\hat{v}$
- Then get **score vector**  $v'$  and softmax it
- Train with cross-entropy

# Word2vec (Mikolov et al.)

## Skip-gram

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

word2vec  
Intuition

Training


Visualization

Semantics  
Miscellany

node2vec

Summary

- Symmetric to CBOW: use  $i$ th row of  $W$  as embedding
- Goal is to maximize  $P(w_{t-2}, w_{t-1}, w_{t+1}, w_{t+2} \mid w_t)$
- Same as minimizing  $-\log P(w_{t-2}, w_{t-1}, w_{t+1}, w_{t+2} \mid w_t)$
- Assume words are independent given  $w_t$ :  
$$P(w_{t-2}, w_{t-1}, w_{t+1}, w_{t+2} \mid w_t) = \prod_{j \in \{-2, -1, 1, 2\}} P(w_{t+j} \mid w_t)$$



# Word2vec (Mikolov et al.)

## Skip-Gram with Negative Sampling (SGNS)

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectors

word2vec

Intuition

Training

Visualization

Semantics

Miscellany

node2vec

Summary

Given target word  $t$

- ① Treat  $t$  and a nearby **context word**  $c$  as a positive example, e.g.,

... lemon, a tablespoon of apricot jam, a pinch ...

$c_1 \quad c_2 \quad t \quad c_3 \quad c_4$

- ② Randomly sample other words for negative examples

- E.g.,  $t = \text{"apricot"}$ ,  $c = \text{"aardvark"}$
- Probability of selecting word  $w$  as a negative is

$$\frac{\text{count}(w)^\alpha}{\sum'_w \text{count}(w')^\alpha}$$

for parameter  $\alpha$  (e.g., 0.75)

- ③ Train with logistic output activation, starting with random weights
- ④ Use final weights as embeddings

# Word2vec (Mikolov et al.)

## Skip-Gram with Negative Sampling (SGNS)

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectors

word2vec

Intuition

Training

Visualization

Semantics

Miscellany

node2vec

Summary

- Given  $t$  and  $c$ , goal is to predict probability that they are near each other (“+” label)
- Represent  $t$  and  $c$  as their embedded representations and compute similarity via dot product:  $t \cdot c$
- Use logistic function to map to a probability in interval  $[0, 1]$ :

$$P(+ | t, c) = \frac{1}{1 + \exp(-t \cdot c)}$$

$$P(- | t, c) = 1 - P(+ | t, c) = \frac{\exp(-t \cdot c)}{1 + \exp(-t \cdot c)}$$

- For  $k$  context words, log of total probability is

$$\log P(+ | t, c_1, \dots, c_k) = \sum_{i=1}^k \log \frac{1}{1 + \exp(-t \cdot c_i)}$$

# Word2vec (Mikolov et al.)

## Skip-Gram with Negative Sampling (SGNS)

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectors

word2vec

Intuition

Training

Visualization

Semantics

Miscellany

node2vec

Summary

- Given positive pairs  $(t, c) \in S_+$  and  $k$  negative pairs  $(t, n) \in S_-$ , want to maximize

$$\begin{aligned} & \sum_{(t,c) \in S_+} \log P(+) \mid t, c) + \sum_{(t,n) \in S_-} \log P(- \mid t, n) \\ = & \sum_{(t,c) \in S_+} \log \sigma(c \cdot t) + \sum_{(t,n) \in S_-} \log \sigma(-n \cdot t) \\ = & \sum_{(t,c) \in S_+} \log \frac{1}{1 + \exp(-c \cdot t)} + \sum_{(t,n) \in S_-} \log \frac{1}{1 + \exp(n \cdot t)} \end{aligned}$$

- I.e., optimizing involves finding an embedding maximizing  $c \cdot t$  and minimizing  $n \cdot t$  for all  $t$

# Word2vec (Mikolov et al.)

## Skip-Gram with Negative Sampling (SGNS)

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

word2vec

Intuition

Training

Visualization


Semantics

Miscellany

node2vec

Summary



- Because word  $w$  can be used both as a target word and a context word, training results in two matrices: **target embedding  $T$**  and **context embedding  $C$**
- $T$  and  $C$  form parameters  $\theta$ 
  - Row  $w$  of  $T$  is  $d$ -dimensional representation of word  $w$
  - Column  $w$  of  $C$  is  $d$ -dim representation of word  $w$
- Often will just use  $T$  as final embedding, but could concatenate  $w$ 's vectors into a  $2d$ -dimensional vector



# Word2vec (Mikolov et al.)

## Visualization of Embeddings

- High-dim representations make visualizing embeddings difficult
  - One visualization technique:  
**Hierarchically cluster** words into a **dendogram**
  - Another common approach:  
Project to 2-dimensional space using **t-SNE**



# Word2vec (Mikolov et al.)

## Semantics

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectors

word2vec

Intuition

Training

Visualization


Semantics

Miscellany

node2vec

Summary

- Embeddings can allow **semantic arithmetic**
- E.g., take embedded representation of “king”, subtract “man” and add “woman”, and result is near that of embedded representation of “queen”
- (“queen” minus “king”)  $\approx$  (“woman” minus “man”)



# Word2vec (Mikolov et al.)

## Semantics

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectors

word2vec

Intuition

Training


Visualization

Semantics

Miscellany

node2vec

Summary




Adjective → comparative → superlative

# Word2vec (Mikolov et al.)

## Semantics

Embeddings  
Stephen Scott  
Introduction  
Vector semantics  
Words and Vectors  
word2vec  
Intuition  
Training  
Visualization  
Semantics  
Miscellany  
node2vec  
Summary



Distances between countries and capitals similar

# Word2vec

## Miscellany

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

word2vec

Intuition

Training

Visualization

Semantics

Miscellany

node2vec

Summary

- Semantics can also reveal biases in corpus
  - E.g., “man” minus “computer programmer” plus “woman” yields “homemaker”
  - Racial biases also revealed
- Application of word2vec to **materials science paper abstracts**
  - Embedded vectors of elements appearing in abstracts clustered according to the periodic table
  - Was able to predict properties of materials before discovered by scientists

# Node2vec (Grover and Leskovec, 2016)

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

word2vec

node2vec

Summary

- Word2vec's approach generalizes beyond text
- All we need to do is represent the context of an instance to embed together instances with similar contexts
  - E.g., biological sequences, nodes in a graph
- Node2vec defines its context for a node based on its local neighborhood, role in the graph, etc.

# Node2vec (Grover and Leskovec, 2016)

## Notation

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

word2vec

node2vec

Summary

- $\mathcal{G} = (\mathcal{V}, \mathcal{E})$
- $\mathcal{A}$  is a  $|\mathcal{V}| \times |\mathcal{V}|$  adjacency matrix
- $f : \mathcal{V} \rightarrow \mathbb{R}^d$  is a mapping function from individual nodes to feature representations
  - $|\mathcal{V}| \times d$  matrix
- $N_S(u) \subset \mathcal{V}$  denotes a neighborhood of node  $u$  generated through a **neighborhood sampling strategy**  $S$
- **Objective:** Preserve local neighborhoods of nodes

# Node2vec (Grover and Leskovec, 2016)

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectors


word2vec

node2vec

Summary

Organization of nodes is based on:

- **Homophily:** Nodes that are highly interconnected and cluster together should embed near each other
- **Structural roles:** Nodes with similar roles in the graph (e.g., hubs) should embed near each other
  - $u$  and  $s_1$  belong to the same community of nodes
  - $u$  and  $s_6$  in two distinct communities share same structural role of a hub node



## Goal

- Embed nodes from the same network community closely together
- Nodes that share similar roles have similar embeddings

# Node2vec

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectors

word2vec

node2vec

Summary


**Key Contribution:** Defining a flexible notion of a node's network neighborhood.

## ① BFS: role of the vertex

- Far apart from each other but share similar kind of vertices

## ② DFS: community

- Reachability/closeness of the two nodes
- My friend's friend's friend has a higher chance to belong to the same community as me



# Node2vec

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectors

word2vec

node2vec

Summary

## Objective function

$$\max_f \sum_{u \in \mathcal{V}} \log P(N_S(u) | f(u))$$

Assumptions:

- Conditional independence:

$$P(N_S(u) | f(u)) = \prod_{n_i \in N_S(u)} P(n_i | f(u))$$

- Symmetry in feature space:

$$P(n_i | f(u)) = \frac{\exp(f(n_i) \cdot f(u))}{\sum_{v \in \mathcal{V}} \exp(f(v) \cdot f(u))}$$

Objective function simplifies to:

$$\max_f \sum_{u \in \mathcal{V}} \left[ -\log Z_u + \sum_{n_i \in N_S(u)} f(n_i) \cdot f(u) \right]$$

# Node2vec (Grover and Leskovec, 2016)

## Neighborhood Sampling

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

word2vec

node2vec

Summary

Given a source node  $u$ , we simulate a random walk of fixed length  $\ell$ :

$$P(c_i = x \mid c_{i-1} = v) = \begin{cases} \frac{\pi_{vx}}{Z} & \text{if } (v, x) \in \mathcal{E} \\ 0 & \text{otherwise} \end{cases}$$

- $c_0 = u$
- $\pi_{vx}$  is the unnormalized transition probability
- $Z$  is the normalization constant.
- 2<sup>nd</sup> order Markovian

# Node2vec (Grover and Leskovec, 2016)

## Neighborhood Sampling

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectors


word2vec

node2vec

Summary

**Search bias**  $\alpha$ :  $\pi_{vx} = \alpha_{pq}(t, x)w_{vx}$  where

$$\alpha_{pq}(t, x) = \begin{cases} \frac{1}{p} & \text{if } d_{tx} = 0 \\ 1 & \text{if } d_{tx} = 1 \\ \frac{1}{q} & \text{if } d_{tx} = 2 \end{cases}$$



**Return parameter**  $p$ :

- Controls the likelihood of immediately revisiting a node in the walk
- If  $p > \max(q, 1)$ 
  - less likely to sample an already visited node
  - avoids 2-hop redundancy in sampling
- If  $p < \min(q, 1)$ 
  - backtrack a step
  - keep the walk local

# Node2vec (Grover and Leskovec, 2016)

## Neighborhood Sampling

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

word2vec

node2vec

Summary

### In-out parameter $q$ :

- If  $q > 1$  inward exploration
  - Local view
  - BFS behavior
- If  $q < 1$  outward exploration
  - Global view
  - DFS behavior

# Node2vec (Grover and Leskovec, 2016)

## Algorithm

Embeddings

Stephen Scott

Introduction

Vector  
semanticsWords and  
Vectors

word2vec

node2vec

Summary

**Algorithm 1** The node2vec algorithm.

---

**LearnFeatures** (Graph  $G = (V, E, W)$ , Dimensions  $d$ , Walks per node  $r$ , Walk length  $l$ , Context size  $k$ , Return  $p$ , In-out  $q$ )
 $\pi = \text{PreprocessModifiedWeights}(G, p, q)$ 
 $G' = (V, E, \pi)$ Initialize  $walks$  to Empty
**for**  $iter = 1$  to  $r$  **do**
     **for all** nodes  $u \in V$  **do**
          $walk = \text{node2vecWalk}(G', u, l)$ 
         Append  $walk$  to  $walks$ 
 $f = \text{StochasticGradientDescent}(k, d, walks)$ 
**return**  $f$ 


---

**node2vecWalk** (Graph  $G' = (V, E, \pi)$ , Start node  $u$ , Length  $l$ )Initialize  $walk$  to  $[u]$ 
**for**  $walk\_iter = 1$  to  $l$  **do**
      $curr = walk[-1]$ 
      $V_{curr} = \text{GetNeighbors}(curr, G')$ 
      $s = \text{AliasSample}(V_{curr}, \pi)$ 
     Append  $s$  to  $walk$ 
**return**  $walk$ 


---

- Implicit bias due to choice of the start node  $u$

- Simulating  $r$  random walks of fixed length  $l$  starting from every node

## Phases:

- 1 Preprocessing to compute transition probabilities
- 2 Random walks
- 3 Optimization using SGD

Each phase is parallelizable and executed asynchronously



# Summary

Embeddings

Stephen Scott

Introduction

Vector  
semantics

Words and  
Vectors

word2vec

node2vec

Summary

- Numeric representations of words called **embeddings** are an important part of NLP
- One type of **sparse** embedding is based on **co-occurrence** counts, often weighted by **tf-idf**
- A popular **dense** embedding is **word2vec**, which is built via training a logistic classifier
- Other popular embeddings are **GloVe** and **fasttext**
- **Dot products** of embedded vectors (which equals cosine when vectors are unit length) are good ways of capturing word similarity
- **Semantic arithmetic** is possible with some embeddings, including adding and subtracting vectors
- The word2vec approach is applicable to other types of data with contexts, such as **nodes in a graph**