

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary

# Regularization and Performance Evaluation

Stephen Scott and Vinod Variyam

[sscott2@unl.edu](mailto:sscott2@unl.edu)

# Introduction

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary

- Machine learning can generally be distilled to an optimization problem
- Choose a classifier (function, hypothesis) from a set of functions that minimizes an objective function
- Clearly we want part of this function to measure performance on the training set, but this is insufficient

# Outline

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary

- Types of machine learning problems
- Loss functions
- Generalization performance vs training set performance
- Overfitting
- Regularization
- Estimating generalization performance
- Summary

# Machine Learning Problems

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary

- **Supervised Learning:** Algorithm given labeled **training data** and infers function (**hypothesis**) from a family of functions (e.g., set of all ANNs) that is able to predict well on new, unseen examples
  - **Classification:** Labels come from a finite, discrete set
  - **Regression:** Labels are real-valued
- **Unsupervised Learning:** Algorithm is given data without labels and is asked to model its structure
  - Clustering, density estimation
- **Reinforcement Learning:** Algorithm controls an agent that interacts with its environment and learns good actions in various situations

# Measuring Performance Loss

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Loss  
Overfitting

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary  
5/54

- In any learning problem, need to be able to quantify performance of algorithm
- In supervised learning, we often use **loss function** (or error function)  $J$  for this task
- Given instance  $x$  with true label  $y$ , if the learner's prediction on  $x$  is  $\hat{y}$ , then

$$J(y, \hat{y})$$

is the loss on that instance

# Measuring Performance

## Examples of Loss Functions

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Loss

Overfitting

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary  
6 / 54

- **0-1 Loss:**  $J(y, \hat{y}) = 1$  if  $y \neq \hat{y}$ , 0 otherwise
- **Square Loss:**  $J(y, \hat{y}) = (y - \hat{y})^2$
- **Cross-Entropy:**  $J(y, \hat{y}) = -y \ln \hat{y} - (1 - y) \ln (1 - \hat{y})$   
( $y$  and  $\hat{y}$  are considered probabilities of a '1' label)
  - Generalizes to  $k$  classes ( $i^*$  = correct class):

$$J(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{i=1}^k y_i \ln \hat{y}_i = - \ln \hat{y}_{i^*}$$

( $\mathbf{y}$  is one-hot vector;  $\hat{y}_i$  is predicted prob. of class  $i$ )

- **Hinge Loss:**  $J(y, \hat{y}) = \max(0, 1 - y\hat{y})$   
(used sometimes for large margin classifiers like SVMs)

All non-negative

# Measuring Performance

## Training Loss

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Loss  
Overfitting

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary  
7 / 54

- Given a loss function  $J$  and a training set  $\mathcal{X}$ , the total loss of the classifier  $h$  on  $\mathcal{X}$  is

$$\text{error}_{\mathcal{X}}(h) = \sum_{x \in \mathcal{X}} J(y_x, \hat{y}_x) ,$$

where  $y_x$  is  $x$ 's label and  $\hat{y}_x$  is  $h$ 's prediction

# Measuring Performance

## Expected Loss

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction  
Measuring  
Performance

Loss  
Overfitting

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary  
8 / 54

- More importantly, the learner needs to **generalize** well:  
Given a new example drawn iid according to unknown probability distribution  $\mathcal{D}$ , we want to minimize  $h$ 's **expected loss**:

$$\text{error}_{\mathcal{D}}(h) = \mathbb{E}_{x \sim \mathcal{D}} [J(y_x, \hat{y}_x)]$$

- Is minimizing training loss the same as minimizing expected loss?

# Measuring Performance

## Expected vs Training Loss

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction  
Measuring  
Performance  
Loss  
Overfitting

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary  
9 / 54

- Sufficiently sophisticated learners (decision trees, multi-layer ANNs) can often achieve arbitrarily small (or zero) loss on a training set
- A hypothesis (e.g., ANN with specific parameters)  $h$  **overfits** the training data  $\mathcal{X}$  if there is an alternative hypothesis  $h'$  such that

$$\text{error}_{\mathcal{X}}(h) < \text{error}_{\mathcal{X}}(h')$$

and

$$\text{error}_{\mathcal{D}}(h) > \text{error}_{\mathcal{D}}(h')$$

# Measuring Performance

## Overfitting

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Loss  
Overfitting

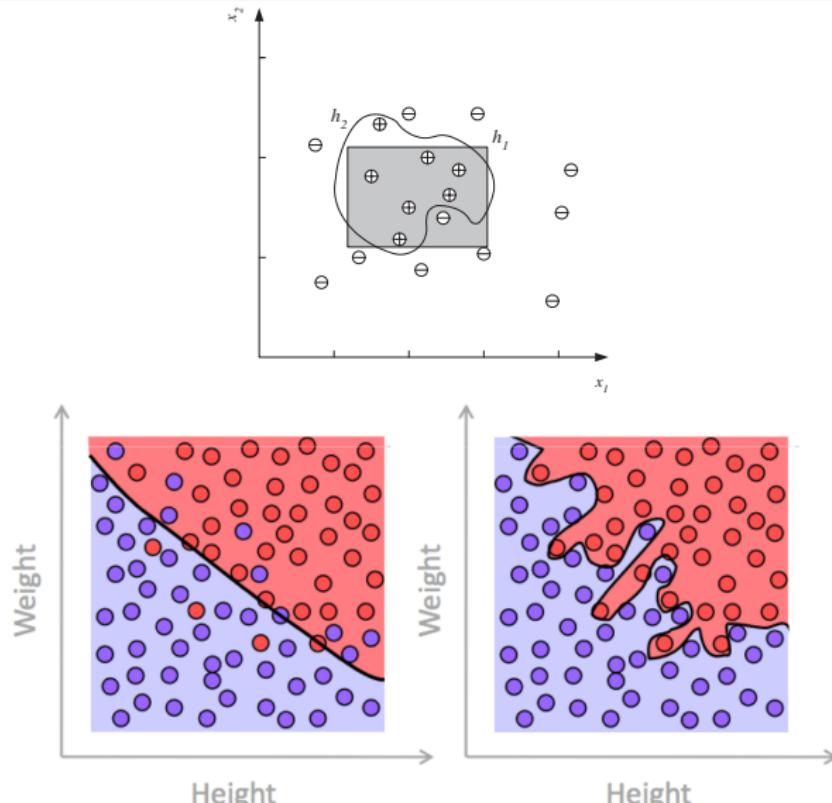
Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary  
10 / 54



# Measuring Performance

## Overfitting

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Loss  
Overfitting

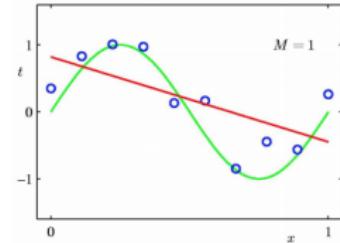
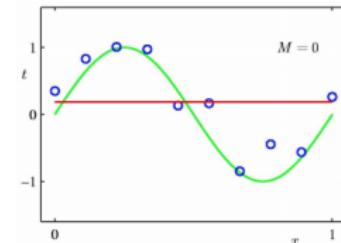
Regularization

Estimating  
Generalization  
Performance

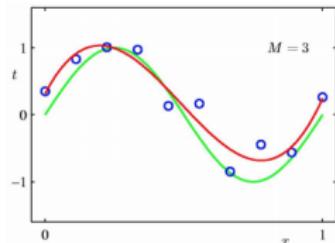
Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

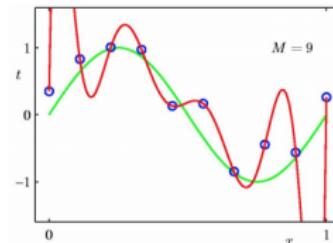
Summary  
11 / 54



Poor  
representations  
of  $\sin(2\pi x)$



Best Fit  
to  
 $\sin(2\pi x)$



Over Fit  
Poor  
representation  
of  $\sin(2\pi x)$

To generalize well, need to balance **training accuracy** with **simplicity**

# Regularization

## Causes of Overfitting

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Causes of Overfitting

Early Stopping

Parameter Norm  
Penalties

Data Augmentation

Multitask Learning

Dropout

Batch Normalization

Others

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

- Generally, if the set of functions  $\mathcal{H}$  the learner has to choose from is complex relative to what is required for correctly predicting the labels of  $\mathcal{X}$ , there's a larger chance of overfitting due to the large number of "wrong" choices in  $\mathcal{H}$ 
  - Could be due to an overly sophisticated set of functions
    - E.g., can fit any set of  $n$  real-valued points with an  $(n - 1)$ -degree polynomial, but perhaps only degree 2 is needed
    - E.g., using an ANN with 5 hidden layers to solve the logical AND problem
  - Could be due to training an ANN too long
    - Over-training an ANN often leads to weights deviating far from zero
    - Makes the function more non-linear, and more complex
  - Often, a larger data set mitigates the problem

# Regularization

## Causes of Overfitting: Overtraining

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Causes of Overfitting

Early Stopping

Parameter Norm  
Penalties

Data Augmentation

Multitask Learning

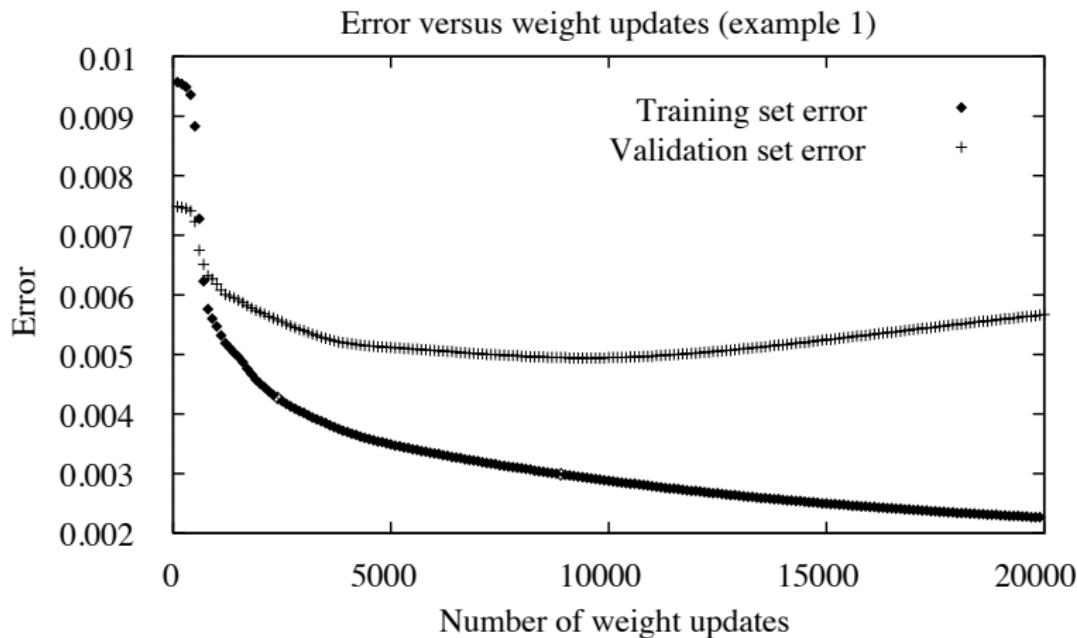
Dropout

Batch Normalization

Others

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms



# Regularization

## Early Stopping

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Causes of Overfitting

Early Stopping

Parameter Norm  
Penalties

Data Augmentation

Multitask Learning

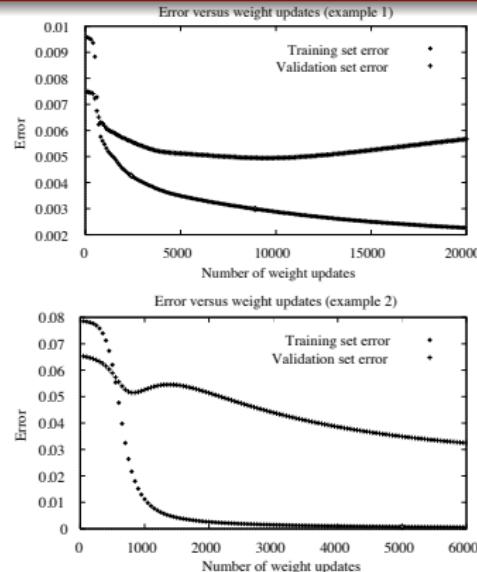
Dropout

Batch Normalization

Others

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms



- Danger of stopping too soon
  - “Patience” parameter determines how long to wait
- Can re-start and track best one on separate validation set

# Regularization

## Parameter Norm Penalties

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction  
Measuring  
Performance

Regularization  
Causes of Overfitting  
Early Stopping  
Parameter Norm  
Penalties

Data Augmentation  
Multitask Learning  
Dropout  
Batch Normalization  
Others

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

- Still want to minimize training loss, but balance it against a **complexity penalty** on the parameters used:

$$\tilde{J}(\boldsymbol{\theta}; \mathcal{X}, \mathbf{y}) = J(\boldsymbol{\theta}; \mathcal{X}, \mathbf{y}) + \alpha \Omega(\boldsymbol{\theta})$$

- $\alpha \in [0, \infty)$  weights loss  $J$  against penalty  $\Omega$

# Regularization

## Parameter Norm Penalties: $L^2$ Norm

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Causes of Overfitting

Early Stopping

Parameter Norm  
Penalties

Data Augmentation

Multitask Learning

Dropout

Batch Normalization

Others

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

- $\Omega(\theta) = (1/2)\|\theta\|_2^2$ , i.e., sum of squares of network's weights
- Since  $\theta = w$ , this becomes

$$\tilde{J}(w; \mathcal{X}, y) = (\alpha/2)w^\top w + J(w; \mathcal{X}, y)$$

- As weights deviate from zero, activation functions become more nonlinear, which is higher risk of overfitting

# Regularization

Parameter Norm Penalties:  $L^2$  Norm

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Causes of Overfitting

Early Stopping

Parameter Norm  
Penalties

Data Augmentation

Multitask Learning

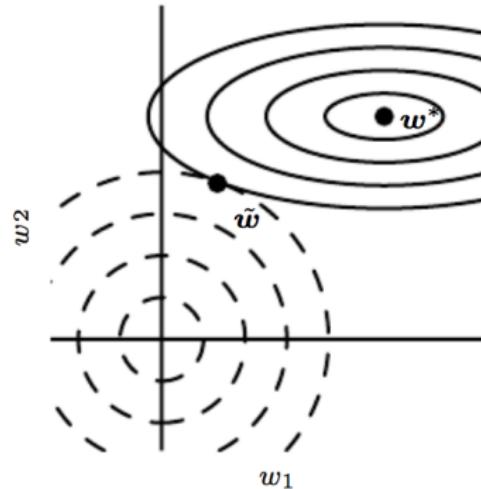
Dropout

Batch Normalization

Others

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms



- $w^*$  is optimal for  $J$ ,  $\mathbf{0}$  optimal for regularizer
- $J$  less sensitive to  $w_1$ , so  $\tilde{w}$  (optimal for  $\tilde{J}$ ) closer to  $w_2$  axis than  $w_1$

# Regularization

## Parameter Norm Penalties: $L^1$ Norm

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization  
Causes of Overfitting  
Early Stopping

Parameter Norm  
Penalties

Data Augmentation  
Multitask Learning  
Dropout  
Batch Normalization  
Others

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

- $\Omega(\theta) = \|\theta\|_1$ , i.e., sum of absolute values of network's weights

$$\tilde{J}(\mathbf{w}; \mathcal{X}, \mathbf{y}) = \alpha \|\mathbf{w}\|_1 + J(\mathbf{w}; \mathcal{X}, \mathbf{y})$$

- As with  $L^2$  regularization, penalizes large weights
- Unlike  $L^2$  regularization, can drive some weights to zero

- **Sparse** solution
- Sometimes used in **feature selection** (e.g., LASSO algorithm)

# Regularization

## Data Augmentation

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Causes of Overfitting

Early Stopping

Parameter Norm

Penalties

Data Augmentation

Multitask Learning

Dropout

Batch Normalization

Others

Estimating  
Generalization  
Performance

Comparing  
Learning

Algorithms

- If  $\mathcal{H}$  powerful and  $\mathcal{X}$  small, then learner can choose some  $h \in \mathcal{H}$  that fits idiosyncrasies or noise in data
- Deep ANNs would like to have at least thousands or tens of thousands of data points
- In classification of high-dimensional data (e.g., image classification), want learned classifier to tolerate transformations and noise
  - ⇒ Can artificially enlarge data set by duplicating existing instances and applying transformations
    - Translating, rotating, scaling
    - Don't change the class, e.g., "b" vs "d" or "6" vs "9"
    - **Don't let duplicates lie in both training and testing sets**
  - ⇒ Can also apply noise injection to input or hidden layers

# Regularization

## Data Augmentation

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

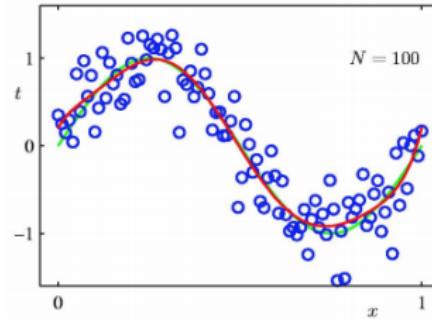
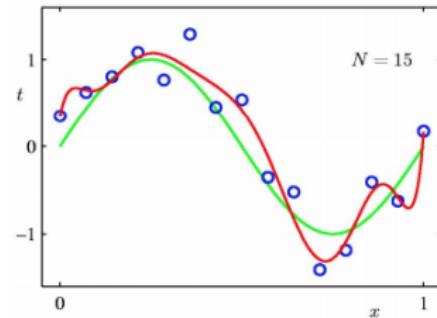
Causes of Overfitting  
Early Stopping  
Parameter Norm  
Penalties

Data Augmentation

Multitask Learning  
Dropout  
Batch Normalization  
Others

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms



# Regularization

## Multitask Learning

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction  
Measuring  
Performance

Regularization  
Causes of Overfitting

Early Stopping  
Parameter Norm  
Penalties

Data Augmentation

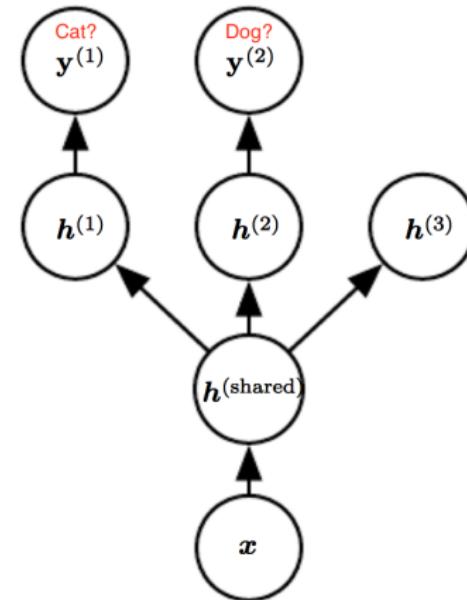
Multitask Learning

Dropout  
Batch Normalization  
Others

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

- If multiple tasks share **generic parameters**, initially process inputs via shared nodes, then do final processing via task-specific nodes
- Backpropagation works as before with multiple output nodes
- Serves as a regularizer since parameter tuning of shared nodes is based on backpropagated error from multiple tasks



# Regularization

## Dropout

Regularization  
and  
Performance  
Evaluation  
Stephen Scott  
and Vinod  
Variyam

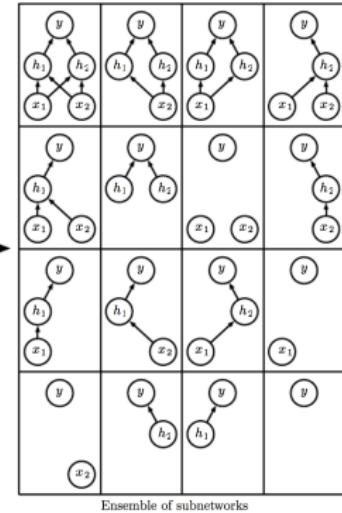
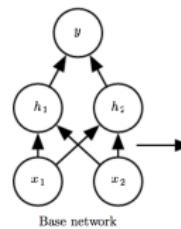
Introduction  
Measuring  
Performance

Regularization  
Causes of Overfitting  
Early Stopping  
Parameter Norm  
Penalties  
Data Augmentation  
Multitask Learning  
Dropout  
Batch Normalization  
Others

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

- Imagine if, for a network, we could average over all networks with each subset of nodes deleted
- Analogous to **bagging**, where we average over ANNs trained on random samples of  $\mathcal{X}$
- In each training iteration, sample a random bit vector  $\mu$ , which determines which nodes are used (e.g.,  $P(\mu_i = 1) = 0.8$  for input unit, 0.5 for hidden unit)



- When training done, re-scale by  $P(\mu_i = 1)$

# Regularization

Batch Normalization (Ioffe and Szegedy 2015)

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Causes of Overfitting

Early Stopping

Parameter Norm  
Penalties

Data Augmentation

Multitask Learning

Dropout

Batch Normalization

Others

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

- Addresses **internal covariate shift**, where changing parameters of layer  $i$  changes distribution of inputs to layer  $i + 1$

- Related to  **$\mathbb{Z}$ -normalization**, where one subtracts sample mean and scales with standard deviation

- $\gamma, \beta$  learnable parameters
- Allows use of higher learning rates, possibly speeding convergence
- In some cases, reduces/eliminates need for dropout

**Input:** Values of  $x$  over a mini-batch:  $B = \{x_1 \dots m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = BN_{\gamma, \beta}(x_i)\}$

$$\mu_B \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_B^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2 \quad // \text{mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv BN_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.

# Regularization

## Other Approaches

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Causes of Overfitting

Early Stopping

Parameter Norm

Penalties

Data Augmentation

Multitask Learning

Dropout

Batch Normalization

Others

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

- **Parameter Tying:** If two learners are learning the same task but different scenarios (distributions, etc.), can tie their parameters together

- If  $w^{(A)}$  are weights for task A and  $w^{(B)}$  are weights for task B, then can use regularization term

$$\Omega(w^{(A)}, w^{(B)}) = \|w^{(A)} - w^{(B)}\|_2^2$$

- E.g., A is supervised and B is unsupervised

- **Parameter Sharing:** When detecting objects in an image, the same recognizer should apply invariant to translation

- Train a single detector (subnetwork) for an object (e.g., cat) by training full network on multiple images with translated cats, where the cat detector subnets share parameters (single copy, used multiple times)

# Regularization

## Other Approaches (cont'd)

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Causes of Overfitting

Early Stopping

Parameter Norm  
Penalties

Data Augmentation

Multitask Learning

Dropout

Batch Normalization

Others

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

- **Sparse Representations:** Instead of penalizing large weights, penalize large outputs of hidden nodes:

$$\tilde{J}(\boldsymbol{\theta}; \mathcal{X}, \mathbf{y}) = J(\boldsymbol{\theta}; \mathcal{X}, \mathbf{y}) + \alpha \Omega(\mathbf{h}) ,$$

where  $\mathbf{h}$  is the vector of hidden unit outputs

# Estimating Generalization Performance

## Setting Goals

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Setting Goals

Confidence Intervals

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary

26 / 54

- Before setting up an experiment, need to understand exactly what the goal is
  - Estimate generalization performance of a hypothesis
  - Tuning hyperparameters
  - Comparing two learning algorithms on a specific task
  - Etc.
- Will never be able to answer the question with 100% certainty
  - Due to variances in training set selection, test set selection, etc.
  - Will choose an **estimator** for the quantity in question, determine the probability distribution of the estimator, and bound the probability that the estimator is way off
  - Estimator needs to work regardless of distribution of training/testing data

# Estimating Generalization Performance

## Setting Goals

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Setting Goals

Confidence Intervals

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary  
27 / 54

- Need to note that, in addition to statistical variations, what we determine is limited to the application that we are studying
  - E.g., if ANN<sub>1</sub> better than ANN<sub>2</sub> on speech recognition, that means nothing about video analysis
- In planning experiments, need to ensure that training data not used for evaluation
  - I.e., **don't test on the training set!**
  - Will bias the performance estimator
  - If using data augmentation, **don't let duplicates lie in both training and testing sets**
  - Also holds for **validation set** used for early stopping, tuning parameters, etc.
    - Validation set serves as part of training set, but not used for model building

# Confidence Intervals

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance  
Setting Goals

Confidence Intervals

Let  $\text{error}_{\mathcal{D}}(h)$  be 0-1 loss of hypothesis  $h$  on instances drawn according to distribution  $\mathcal{D}$ . If

- Test set  $\mathcal{V}$  contains  $N$  examples, drawn independently of  $h$  and each other
- $N \geq 30$

Then with approximately 95% probability,  $\text{error}_{\mathcal{D}}(h)$  lies in

$$\text{error}_{\mathcal{V}}(h) \pm 1.96 \sqrt{\frac{\text{error}_{\mathcal{V}}(h)(1 - \text{error}_{\mathcal{V}}(h))}{N}}$$

E.g., hypothesis  $h$  misclassifies 12 of the 40 examples in test set  $\mathcal{V}$ :

$$\text{error}_{\mathcal{V}}(h) = \frac{12}{40} = 0.30$$

Then with approx. 95% confidence,  $\text{error}_{\mathcal{D}}(h) \in [0.158, 0.442]$

# Confidence Intervals (cont'd)

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction  
Measuring  
Performance  
Regularization  
Estimating  
Generalization  
Performance  
Setting Goals  
Confidence Intervals

Comparing  
Learning  
Algorithms  
Other  
Performance  
Measures

Summary  
29 / 54

Let  $\text{error}_{\mathcal{D}}(h)$  be 0-1 loss of  $h$  on instances drawn according to distribution  $\mathcal{D}$ . If

- Test set  $\mathcal{V}$  contains  $N$  examples, drawn independently of  $h$  and each other
- $N \geq 30$

Then with approximately c% probability,  $\text{error}_{\mathcal{D}}(h)$  lies in

$$\text{error}_{\mathcal{V}}(h) \pm z_c \sqrt{\frac{\text{error}_{\mathcal{V}}(h)(1 - \text{error}_{\mathcal{V}}(h))}{N}}$$

$N\%$ :	50%	68%	80%	90%	95%	98%	99%
$z_c$ :	0.67	1.00	1.28	1.64	1.96	2.33	2.58

Why?

# error<sub>D</sub>(h) is a Random Variable

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction  
Measuring  
Performance  
Regularization  
Estimating  
Generalization  
Performance  
Setting Goals  
Confidence Intervals

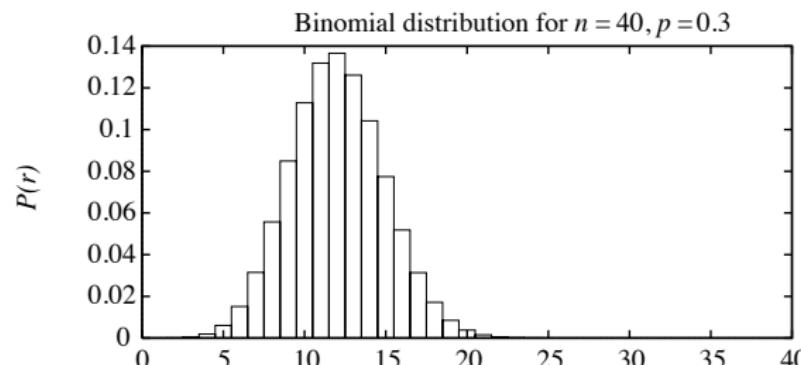
Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary  
30 / 54

Repeatedly run the experiment, each with different randomly drawn  $\mathcal{V}$  (each of size  $N$ )

Probability of observing  $r$  misclassified examples:



$$P(r) = \binom{N}{r} \text{error}_{\mathcal{D}}(h)^r (1 - \text{error}_{\mathcal{D}}(h))^{N-r}$$

I.e., let  $\text{error}_{\mathcal{D}}(h)$  be probability of heads in biased coin, then  
 $P(r) = \text{prob. of getting } r \text{ heads out of } N \text{ flips}$

# Binomial Probability Distribution

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction  
Measuring  
Performance

Regularization  
Estimating  
Generalization  
Performance  
Setting Goals

Confidence Intervals

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary  
31 / 54

$$P(r) = \binom{N}{r} p^r (1-p)^{N-r} = \frac{N!}{r!(N-r)!} p^r (1-p)^{N-r}$$

Probability  $P(r)$  of  $r$  heads in  $N$  coin flips, if  $p = \Pr(\text{heads})$

- Expected, or mean value of  $X$ ,  $\mathbb{E}[X]$  ( $=$  # heads on  $N$  flips  $=$  # mistakes on  $N$  test exs), is

$$\mathbb{E}[X] \equiv \sum_{i=0}^N i P(i) = Np = N \cdot \text{error}_{\mathcal{D}}(h)$$

- Variance of  $X$  is

$$\text{Var}(X) \equiv \mathbb{E}[(X - \mathbb{E}[X])^2] = Np(1-p)$$

- Standard deviation of  $X$ ,  $\sigma_X$ , is

$$\sigma_X \equiv \sqrt{\mathbb{E}[(X - \mathbb{E}[X])^2]} = \sqrt{Np(1-p)}$$

# Approximate Binomial Dist. with Normal

Regularization  
and  
Performance

Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance  
Setting Goals

Confidence Intervals

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary  
32 / 54

$\text{error}_{\mathcal{V}}(h) = r/N$  is binomially distributed, with

- mean  $\mu_{\text{error}_{\mathcal{V}}(h)} = \text{error}_{\mathcal{D}}(h)$  (**unbiased estimator**)
- standard deviation  $\sigma_{\text{error}_{\mathcal{V}}(h)}$

$$\sigma_{\text{error}_{\mathcal{V}}(h)} = \sqrt{\frac{\text{error}_{\mathcal{D}}(h)(1 - \text{error}_{\mathcal{D}}(h))}{N}}$$

(increasing  $N$  decreases variance)

Want to compute confidence interval = interval centered at  $\text{error}_{\mathcal{D}}(h)$  containing  $c\%$  of the weight under the distribution

Approximate binomial by **normal** (Gaussian) dist:

- mean  $\mu_{\text{error}_{\mathcal{V}}(h)} = \text{error}_{\mathcal{D}}(h)$
- standard deviation  $\sigma_{\text{error}_{\mathcal{V}}(h)}$

$$\sigma_{\text{error}_{\mathcal{V}}(h)} \approx \sqrt{\frac{\text{error}_{\mathcal{V}}(h)(1 - \text{error}_{\mathcal{V}}(h))}{N}}$$

# Normal Probability Distribution

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Setting Goals

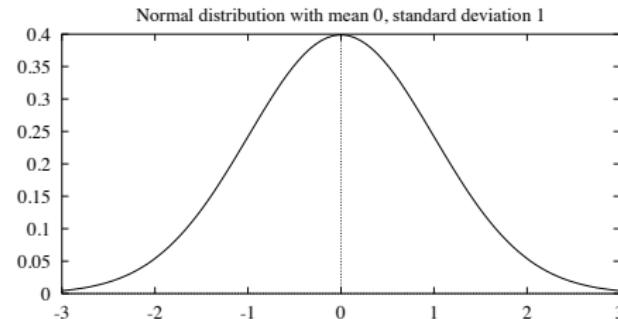
Confidence Intervals

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary

33 / 54



$$p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2\right)$$

- The probability that  $X$  will fall into the interval  $(a, b)$  is given by  $\int_a^b p(x) dx$
- Expected, or mean value of  $X$ ,  $\mathbb{E}[X]$ , is  $\mathbb{E}[X] = \mu$
- Variance is  $Var(X) = \sigma^2$ , standard deviation is  $\sigma_X = \sigma$

# Normal Probability Distribution (cont'd)

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

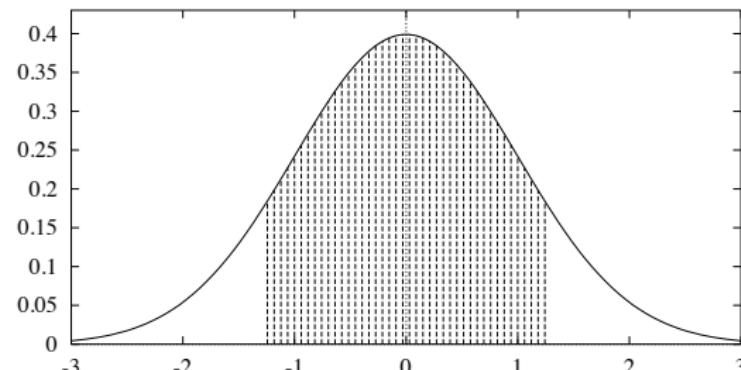
Setting Goals  
Confidence Intervals

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary

34 / 54



80% of area (probability) lies in  $\mu \pm 1.28\sigma$

$c\%$  of area (probability) lies in  $\mu \pm z_c \sigma$

$c\%:$	50%	68%	80%	90%	95%	98%	99%
$z_c:$	0.67	1.00	1.28	1.64	1.96	2.33	2.58

# Normal Probability Distribution (cont'd)

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

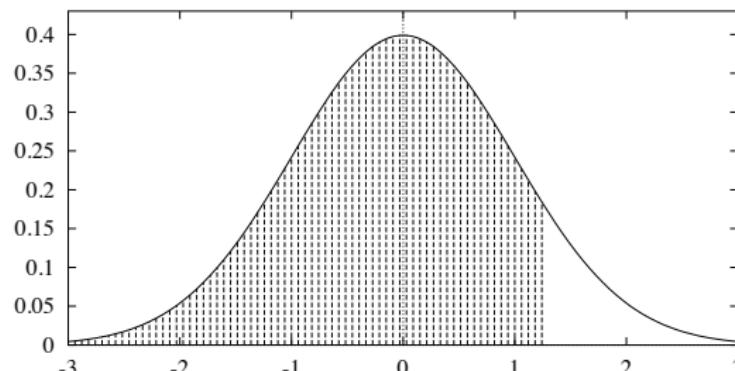
Introduction  
Measuring  
Performance  
Regularization  
Estimating  
Generalization  
Performance  
Setting Goals  
Confidence Intervals

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary  
35 / 54

Can also have **one-sided** bounds:



$c\%$  of area lies  $< \mu + z_c' \sigma$  or  $> \mu - z_c' \sigma$ , where  
 $z_c' = z_{100-(100-c)/2}$

$c\%:$	50%	68%	80%	90%	95%	98%	99%
$z_c':$	0.0	0.47	0.84	1.28	1.64	2.05	2.33

# Confidence Intervals Revisited

Regularization  
and  
Performance

Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Setting Goals

Confidence Intervals

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary  
36 / 54

If  $\mathcal{V}$  contains  $N \geq 30$  examples, indep. of  $h$  and each other

Then with approximately 95% probability,  $error_{\mathcal{V}}(h)$  lies in

$$error_{\mathcal{D}}(h) \pm 1.96 \sqrt{\frac{error_{\mathcal{D}}(h)(1 - error_{\mathcal{D}}(h))}{N}}$$

Equivalently,  $error_{\mathcal{D}}(h)$  lies in

$$error_{\mathcal{V}}(h) \pm 1.96 \sqrt{\frac{error_{\mathcal{D}}(h)(1 - error_{\mathcal{D}}(h))}{N}}$$

which is approximately

$$error_{\mathcal{V}}(h) \pm 1.96 \sqrt{\frac{error_{\mathcal{V}}(h)(1 - error_{\mathcal{V}}(h))}{N}}$$

(One-sided bounds yield upper or lower error bounds)

# Central Limit Theorem

Regularization  
and  
Performance

Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Setting Goals

Confidence Intervals

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary

37 / 54

## How can we justify approximation?

Consider set of iid random variables  $Y_1, \dots, Y_N$ , all from **arbitrary** probability distribution with mean  $\mu$  and finite variance  $\sigma^2$ . Define sample mean  $\bar{Y} \equiv (1/N) \sum_{i=1}^n Y_i$

$\bar{Y}$  is itself a random variable, i.e., result of an experiment (e.g.,  $error_{\mathcal{V}}(h) = r/N$ )

**Central Limit Theorem:** As  $N \rightarrow \infty$ , the distribution governing  $\bar{Y}$  approaches normal distribution with mean  $\mu$  and variance  $\sigma^2/N$

Thus the distribution of  $error_{\mathcal{V}}(h)$  is approximately normal for large  $N$ , and its expected value is  $error_{\mathcal{D}}(h)$

**(Guideline:**  $N \geq 30$  when estimator's distribution is binomial; might need to be larger for other distributions)

# Calculating Confidence Intervals

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction  
Measuring  
Performance  
Regularization  
Estimating  
Generalization  
Performance  
Setting Goals  
**Confidence Intervals**

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary  
38 / 54

- ➊ Pick parameter to estimate:  $\text{error}_{\mathcal{D}}(h)$   
(0-1 loss on distribution  $\mathcal{D}$ )
- ➋ Choose an estimator:  $\text{error}_{\mathcal{V}}(h)$   
(0-1 loss on independent test set  $\mathcal{V}$ )
- ➌ Determine probability distribution that governs estimator:  $\text{error}_{\mathcal{V}}(h)$  governed by binomial distribution, approximated by normal when  $N \geq 30$
- ➍ Find interval  $(L, U)$  such that  $c\%$  of probability mass falls in the interval
  - Could have  $L = -\infty$  or  $U = \infty$
  - Use table of  $z_c$  or  $z'_c$  values (if distribution normal)

# Comparing Learning Algorithms

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

*K*-Fold CV  
Student's *t*  
Distribution

Other  
Performance  
Measures

Summary

- What if we want to compare two learning algorithms  $L^1$  and  $L^2$  (e.g., two ANN architectures, two regularizers, etc.) on a specific application?
- Insufficient to simply compare error rates on a single test set
- Use ***K*-fold cross-validation** and a **paired *t* test**
- Also helpful for **hyperparameter tuning**

# K-Fold Cross Validation

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction  
Measuring  
Performance

Regularization  
Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

K-Fold CV  
Student's  $t$   
Distribution

Other  
Performance  
Measures

Summary

- ➊ Partition data set  $\mathcal{X}$  into  $K$  equal-sized subsets  $\mathcal{X}_1, \mathcal{X}_2, \dots, \mathcal{X}_K$ , where  $|\mathcal{X}_i| \geq 30$
- ➋ For  $i$  from 1 to  $K$ , do  
(Use  $\mathcal{X}_i$  for testing, and rest for training)
  - ➌  $\mathcal{V}_i = \mathcal{X}_i$
  - ➍  $\mathcal{T}_i = \mathcal{X} \setminus \mathcal{X}_i$
  - ➎ Train learning algorithm  $L^1$  on  $\mathcal{T}_i$  to get  $h_i^1$
  - ➏ Train learning algorithm  $L^2$  on  $\mathcal{T}_i$  to get  $h_i^2$
  - ➐ Let  $p_i^j$  be error of  $h_i^j$  on test set  $\mathcal{V}_i$
  - ➑  $p_i = p_i^1 - p_i^2$
- ➒ Error difference estimate  $p = (1/K) \sum_i^K p_i$

# K-Fold Cross Validation (cont'd)

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

K-Fold CV

Student's *t*  
Distribution

Other  
Performance  
Measures

Summary

- Now estimate confidence that true expected error difference  $< 0$
- ⇒ Confidence that  $L^1$  is better than  $L^2$  on learning task
- Use **one-sided test**, with confidence derived from **student's *t* distribution** with  $K - 1$  **degrees of freedom**
- With approximately  $c\%$  probability, true difference of expected error between  $L^1$  and  $L^2$  is at most

$$p + t_{c,K-1} s_p$$

where

$$s_p \equiv \sqrt{\frac{1}{K(K-1)} \sum_{i=1}^K (p_i - p)^2}$$

# Student's $t$ Distribution (One-Sided Test)

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

$K$ -Fold CV

Student's  $t$   
Distribution

Other  
Performance  
Measures

Summary

df	0.600	0.700	0.800	0.900	0.950	0.975	0.990	0.995
1	0.325	0.727	1.376	3.078	6.314	12.706	31.821	63.657
2	0.289	0.617	1.061	1.886	2.920	4.303	6.965	9.925
3	0.277	0.584	0.978	1.638	2.353	3.182	4.541	5.841
4	0.271	0.569	0.941	1.533	2.132	2.776	3.747	4.604
5	0.267	0.559	0.920	1.476	2.015	2.571	3.365	4.032
6	0.265	0.553	0.906	1.440	1.943	2.447	3.143	3.707
7	0.263	0.549	0.896	1.415	1.895	2.365	2.998	3.499
8	0.262	0.546	0.889	1.397	1.860	2.306	2.896	3.355
9	0.261	0.543	0.883	1.383	1.833	2.262	2.821	3.250
10	0.260	0.542	0.879	1.372	1.812	2.228	2.764	3.169
11	0.260	0.540	0.876	1.363	1.796	2.201	2.718	3.106
12	0.259	0.539	0.873	1.356	1.782	2.179	2.681	3.055
13	0.259	0.538	0.870	1.350	1.771	2.160	2.650	3.012

If  $p + t_{c,K-1} s_p < 0$  our assertion that  $L^1$  has less error than  $L^2$  is supported with confidence  $c$

So if  $K$ -fold CV used, compute  $p$ , look up  $t_{c,K-1}$  and check if  $p < -t_{c,K-1} s_p$

**One-sided test; says nothing about  $L^2$  over  $L^1$**

# Caveat

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

K-Fold CV

Student's *t*  
Distribution

Other  
Performance  
Measures

Summary

- Say you want to show that learning algorithm  $L^1$  performs better than algorithms  $L^2, L^3, L^4, L^5$
- If you use  $K$ -fold CV to show superior performance of  $L^1$  over each of  $L^2, \dots, L^5$  at 95% confidence, there's a 5% chance each one is wrong
  - ⇒ There's an over 18.5% chance that at least one is wrong
  - ⇒ Our overall confidence is only just over 81%
- Need to account for this, or use more appropriate test

# More Specific Performance Measures

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Confusion Matrices

ROC Curves

Precision-Recall  
Curves

44 / 54

Summary

- So far, we've looked at a single error rate to compare hypotheses/learning algorithms/etc.
- This may not tell the whole story:
  - 1000 test examples: 20 positive, 980 negative
  - $h^1$  gets 2/20 pos correct, 965/980 neg correct, for accuracy of  $(2 + 965)/(20 + 980) = 0.967$
  - Pretty impressive, except that always predicting negative yields accuracy = 0.980
  - Would we rather have  $h^2$ , which gets 19/20 pos correct and 930/980 neg, for accuracy = 0.949?
  - Depends on how important the positives are, i.e., frequency in practice and/or cost (e.g., cancer diagnosis)

# Confusion Matrices

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Confusion Matrices

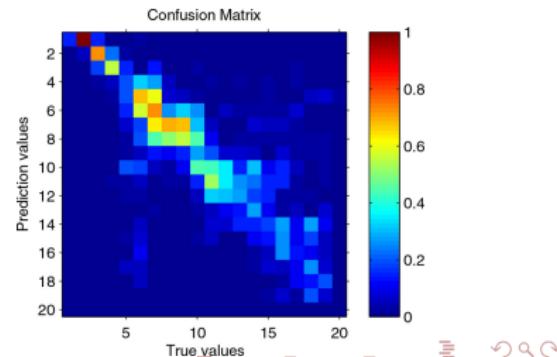
ROC Curves  
Precision-Recall  
Curves

45 / 54  
Summary

Break down error into type: true positive, etc.

		Predicted Class		
True Class		Positive	Negative	Total
Positive	$tp$ : true positive	$fn$ : false negative	$p$	
Negative	$fp$ : false positive	$tn$ : true negative		$n$
Total	$p'$	$n'$		$N$

- Generalizes to multiple classes
- Allows one to quickly assess which classes are missed the most, and into what other class



# ROC Curves

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Confusion Matrices

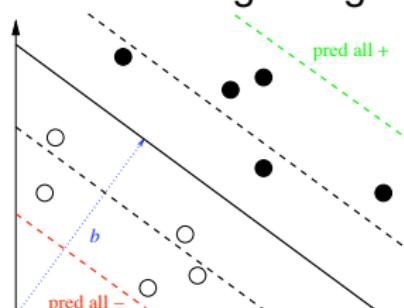
ROC Curves

Precision-Recall  
Curves

46 / 54

Summary

- Consider classification via ANN + linear threshold unit
- Normally threshold  $f(\mathbf{x}; \mathbf{w}, b)$  at 0, but what if we changed it?
- Keeping  $\mathbf{w}$  fixed while changing threshold = fixing hyperplane's slope while moving along its normal vector



- Get a **set** of classifiers, one per labeling of test set
- Similar situation with any classifier with confidence value, e.g., probability-based

# ROC Curves

## Plotting $tp$ versus $fp$

Regularization  
and  
Performance

Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Confusion Matrices

ROC Curves

Precision-Recall  
Curves

47 / 54

Summary

- Consider the “always –” hyp. What is  $fp$ ? What is  $tp$ ? What about the “always +” hyp?
- In between the extremes, we plot TP versus FP by sorting the test examples by the confidence values

Ex	Confidence	label	Ex	Confidence	label
$x_1$	169.752	+	$x_6$	-12.640	-
$x_2$	109.200	+	$x_7$	-29.124	-
$x_3$	19.210	-	$x_8$	-83.222	-
$x_4$	1.905	+	$x_9$	-91.554	+
$x_5$	-2.75	+	$x_{10}$	-128.212	-

# ROC Curves

## Plotting $tp$ versus $fp$ (cont'd)

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

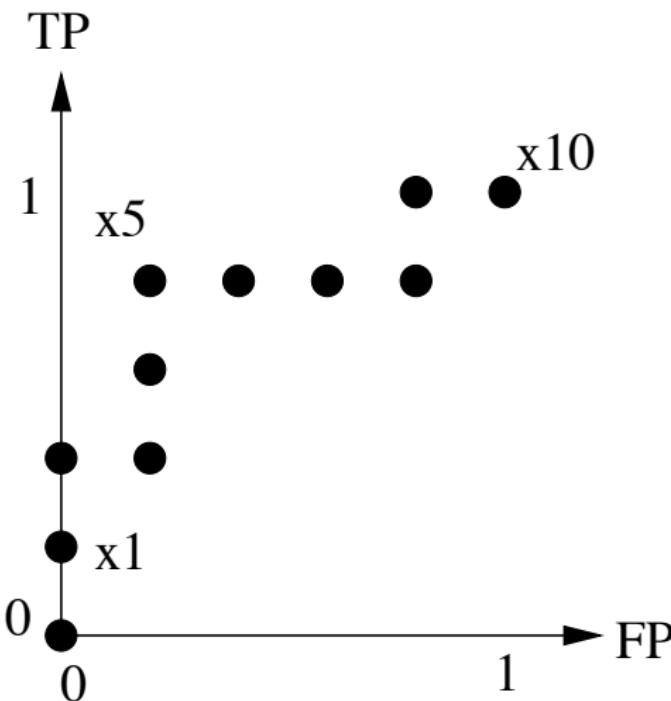
Confusion Matrices

ROC Curves

Precision-Recall  
Curves

48 / 54

Summary



# ROC Curves

## Convex Hull

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

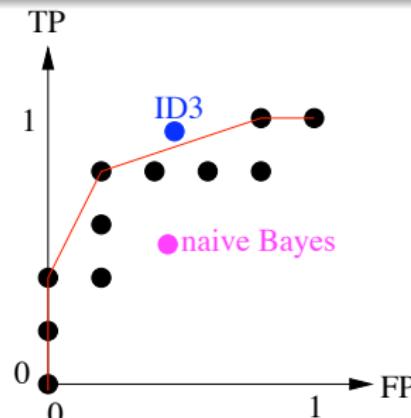
Confusion Matrices

ROC Curves

Precision-Recall  
Curves

49 / 54

Summary



- The **convex hull** of the ROC curve yields a collection of classifiers, each optimal under different conditions
  - If  $FP$  cost =  $FN$  cost, then draw a line with slope  $|N|/|P|$  at  $(0, 1)$  and drag it towards convex hull until you touch it; that's your **operating point**
  - Can use as a classifier any part of the hull since can randomly select between two classifiers

# ROC Curves

## Convex Hull

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

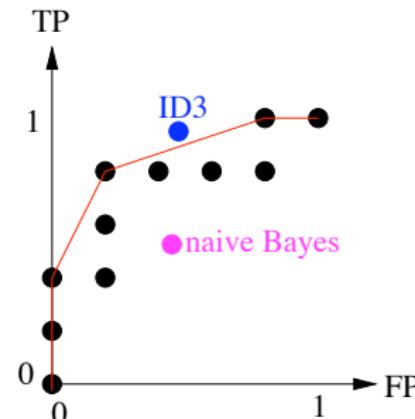
Confusion Matrices

ROC Curves

Precision-Recall  
Curves

50 / 54

Summary



- Can also compare curves against “single-point” classifiers when no curves
  - In plot, ID3 better than our SVM iff negatives scarce; nB never better

# ROC Curves

## Miscellany

Regularization  
and  
Performance  
Evaluation  
  
Stephen Scott  
and Vinod  
Variyam

Introduction  
  
Measuring  
Performance  
  
Regularization  
  
Estimating  
Generalization  
Performance  
  
Comparing  
Learning  
Algorithms  
  
Other  
Performance  
Measures  
  
Confusion Matrices  
ROC Curves

Precision-Recall  
Curves  
51 / 54  
Summary

- What is the worst possible ROC curve?
- One metric for measuring a curve's goodness: **area under curve (AUC)**:

$$\frac{\sum_{x_+ \in P} \sum_{x_- \in N} I(h(x_+) > h(x_-))}{|P| |N|}$$

i.e., rank all examples by confidence in “+” prediction,  
count the number of times a positively-labeled example  
(from  $P$ ) is ranked above a negatively-labeled one (from  
 $N$ ), then normalize

- What is the best value?
- Distribution approximately normal if  $|P|, |N| > 10$ , so can find confidence intervals
- Catching on as a better scalar measure of performance than error rate
- Possible (though tricky) with multi-class problems

# Precision-Recall Curves

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

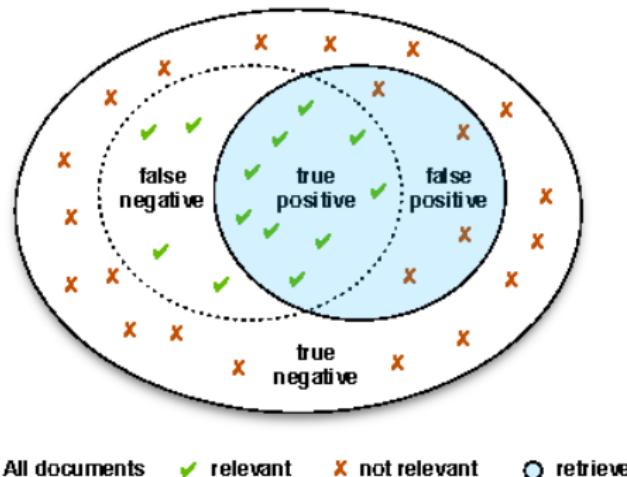
Introduction  
Measuring  
Performance  
Regularization  
Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures  
Confusion Matrices  
ROC Curves

Precision-Recall  
Curves  
52 / 54  
Summary

Consider information retrieval task, e.g., web search



**precision** =  $tp/p' = \text{fraction of retrieved that are positive}$

**recall** =  $tp/p = \text{fraction of positives retrieved}$

# Precision-Recall Curves (cont'd)

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Confusion Matrices

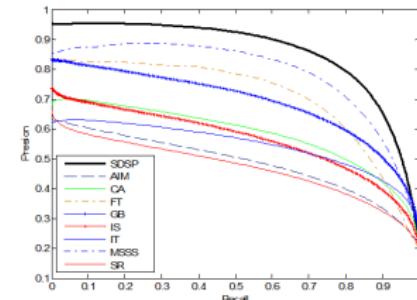
ROC Curves

Precision-Recall  
Curves

53 / 54

Summary

- As with ROC, vary threshold to trade precision and recall
- Can compare curves based on containment
- More suitable than ROC for large numbers of negatives
- Use  $F_\beta$ -measure to combine at a specific point, where  $\beta$  weights precision vs recall:



$$F_\beta \equiv (1 + \beta^2) \frac{precision \cdot recall}{(\beta^2 \cdot precision) + recall}$$

# Summary

Regularization  
and  
Performance  
Evaluation

Stephen Scott  
and Vinod  
Variyam

Introduction

Measuring  
Performance

Regularization

Estimating  
Generalization  
Performance

Comparing  
Learning  
Algorithms

Other  
Performance  
Measures

Summary

- In machine learning, tend to search for a hypothesis that optimizes some **objective function**
- Often try to minimize **loss** on training set, but doing that too aggressively can cause **overfitting**
- Prevent overfitting via **regularization** (e.g., early stopping, complexity penalties, dropout, batch norm)
- Once hypothesis is trained, need to estimate its **generalization performance** by evaluating on an **independent test set**
- Statistical analyses on the test results yields **confidence intervals** and **paired  $t$  tests**
- Other ways of analyzing performance are **confusion matrices**, **ROC curves**, and **precision-recall curves**
- And remember: **Don't test on the training set!**